

# Citation Integrity Audit

Audit of the principal claims and sources used in the Awesome Grounded vs. Ungrounded LLMs repository. The uploaded research paper is the primary organizing source; external links identify the cited papers, official repositories, and benchmark resources.

| Claim / Resource                         | Primary Source        | Verification |
|------------------------------------------|-----------------------|--------------|
| RAG improves factual generation          | Lewis et al. (2020)   | Verified     |
| Retrieval benefits long-tail knowledge   | Mallen et al. (2023)  | Verified     |
| Ungrounded legal hallucination: 58–88%   | Dahl et al. (2024)    | Verified     |
| Grounded legal hallucination: 17–33%     | Magesh et al. (2025)  | Verified     |
| Self-RAG adaptive retrieval and critique | Asai et al. (2024)    | Verified     |
| Atomic factuality evaluation             | Min et al. (2023)     | Verified     |
| Hallucination taxonomy                   | Ji et al. (2023)      | Verified     |
| Truthfulness benchmark                   | Lin et al. (2022)     | Verified     |
| Medical-domain evaluation                | Singhal et al. (2023) | Verified     |

Important limitation: the repository does not claim that grounding makes an LLM hallucination-free. The underlying paper explicitly concludes that grounding improves accuracy but does not guarantee accuracy, and recommends provenance display, human verification, and independent auditing.